

Zero-Day Detection and Mitigation Framework for Kubernetes Microservices: A Temporal Graph Neural Network Approach with Explainable AI

Tejashwini S Uranakar
Information Science and
Engineering

BMS Institute of Technology and
Management
Bengaluru, India
tejashwinisuranakar@ieee.org

Arun Kumar B. R
Information Science and
Engineering

BMS Institute of Technology and
Management
Bengaluru, India
arunkumarbr@bmsit.in

Umar Ismail Shariff
Information Science and
Engineering

BMS Institute of Technology and
Management
Bengaluru, India
1by23is243@bmsit.in

Vineeth Sagar HL
Information Science and Engineering
BMS Institute of Technology and Management
Bengaluru, India
1by23is250@bmsit.in

Vinith S Kiran.
Information Science and Engineering
BMS Institute of Technology and Management
Bengaluru, India
1by23is251@bmsit.in

Abstract— Zero-day exploits pose significant risks to Kubernetes-based microservices due to their dynamic and distributed nature. This paper presents a proof-of-concept framework for detecting and mitigating such threats using temporal graph construction, machine learning, and explainable AI. Multi-source telemetry, including eBPF syscall traces, network flows, service-mesh spans, and Kubernetes audit logs, is aggregated into time-windowed behavioral graphs. An ensemble of baseline anomaly detectors and a Temporal Graph Neural Network (TGNN) enables effective temporal-relational analysis. Explainable AI using SHAP and GNNExplainer maps detected anomalies to MITRE ATT&CK techniques, supporting interpretable alerts. A Kubernetes operator enforces safe containment with human-in-the-loop approvals. Evaluation on a Minikube testbed with synthetic zero-day scenarios achieves 93–96% recall, 85–92% precision ($F1 = 0.94$), and sub-200 ms latency, with ANOVA confirming statistical significance ($p < 0.05$). The framework demonstrates a 15% precision improvement over baseline methods and addresses key gaps in Kubernetes-specific zero-day defense.

Keywords: zero-day detection, temporal graphs, graph neural networks, explainable AI, Kubernetes security, anomaly detection, eBPF, microservices

I. INTRODUCTION

A. Background and Motivation

Cloud-native applications increasingly rely on Kubernetes-based microservices to support scalability and rapid deployment. While this architecture improves flexibility, it also introduces significant security challenges due to highly dynamic workloads, short-lived containers, and complex inter-service communication [1]. These characteristics expand the attack surface and make consistent security monitoring difficult.

Zero-day attacks are particularly dangerous in Kubernetes environments, as they exploit unknown vulnerabilities that cannot be detected using signature-based approaches [2]. Traditional host-based and network-based intrusion detection systems, originally designed for static systems, struggle to

capture evolving behavioral patterns and relationships across microservices. This limitation highlights the need for security mechanisms that model temporal and relational behavior while remaining interpretable and operationally safe.

B. Problem Statement

The existing Kubernetes security mechanisms have a heavy reliance on statistical analysis, policy enforcement, and signature detection methods to maintain safe deployments [2]. The statistical methods are often plagued by a high level of false positive alerts when detecting anomalies, while the signature detection methods do not detect previously unknown zero-day vulnerabilities. Policy-based methods of prevention are not adaptable, and do not provide the user with any adaptive intelligence as it relates to the patterns of a threat. In addition, automated processes without sufficient human input create additional bottlenecks in terms of solution quality and prevent the user from receiving reliable and explainable solutions [11].

C. Proposed Solution

We propose an approach that uses Temporal Graph Neural Networks to combine data from different sources for effective anomaly detection. An Explainable AI layer interprets alerts and helps ensure safe, automated containment within Kubernetes. This design improves real-time detection and response while reducing the impact on service availability [1][2].

Accordingly, this work investigates the following research question: how can temporal-relational modeling combined with explainable AI enable accurate and operationally safe zero-day attack detection and mitigation in Kubernetes-based microservice environments?

D. Contributions

A Kubernetes-aware zero-day attack defense system incorporating multi-source Telemetry, Temporal Graph construction, TGNN anomaly detection, AI Explainability with MITRE ATT&CK Mapping, and safety-oriented Human-in-the-Loop containment is proposed. A proof-of-

concept implementation on Minikube shows a 15% relative accuracy improvement over baseline models.

II. LITERATURE SURVEY AND RELATED WORKS

A. Cloud-Native and Kubernetes Security

Several recent studies have pointed out that there are many security challenges due to the dynamic and distributed nature of containerized workloads on cloud native infrastructure. Arif et al. [1] provided a detailed overview of current privacy-enhancing and trust-centric mechanisms that can be used for protecting data and systems in cloud-native systems, as well as an outline of some of the key architectural principles needed for ensuring data confidentiality and integrity for cloud-native systems. However, the authors did not provide either deployable implementations or Kubernetes-Specific implementations through their research efforts. Additionally, Pitkar [2] examined the use of artificial intelligence (AI) in anomaly detection on Kubernetes workloads, showing how machine learning techniques could provide security improvements for Kubernetes clusters, while also identifying limitations associated with scalability, workload variability, and complex microservice-to-microservice interactions.

B. eBPF-Based Security and Observability

eBPF-based kernel-level observability has become a strong runtime monitoring paradigm for Kubernetes. Hussain et al. [3] prove the efficiency of detection with eBPF-based inspection on attacks such as IP spoofing due to fine-grained syscall visibility. Murugan [4] demonstrates that eBPF enriches visibility at the container level but states "lacking integration at higher levels of machine learning frameworks." ML pipelines on logs, such as studied by Mohammed [5], reach improved anomaly detection from Kubernetes logs, but with weak multi-source fusion.

C. Graph and Temporal Models for Intrusion Detection

Models based on Graphs demonstrate a high level of efficacy at capturing complex interaction dynamics in systems. For instance, Ben-Yair et al. [6] were able to combine eBPF telemetry data with meta learning algorithms to detect anomalies but did not develop temporal or relational models. Fu et al. [9] used self-supervised learning methods on dynamic graphs and found that the results captured long-range structural relationships (temporal and spatial) which enabled better anomaly detection than methods which did not use these types of graph structures. GUIDE-GNN [8] applied GNN techniques in a microservice environment; however, it did not account for temporal properties. More recent development of Transformer-enhanced graph-based intrusion detection models also support the utility of graphs in modelling complex attack patterns [16]. However, for the most part, all of these methods have not been specifically developed to address issues associated with Kubernetes-native telemetry data nor do they offer any asynchronous operational safety features (e.g. explainability and 'human-in-the-loop' containment).

D. Explainable AI in Cybersecurity

Explainability has become crucial for operational trust in security systems. The SETS cybersecurity survey [10] reviews existing XAI techniques and points out that XAI

research lacks standardized evaluation in real-world security deployments. Alenezi and Ludwig [11] demonstrate using SHAP to improve interpretability in cybersecurity datasets; however, this work focuses on feature-level explanations and considers neither graph-structured reasoning. Most prior XAI methods remain post-hoc and loosely coupled to detection pipelines, which motivates tight integration of explainability with graph-based anomaly detection and mitigation workflows in cloud-native environments.

E. Zero-Day and Sequence-Based Detection Approaches

Researchers have looked at anomaly detection techniques like isolation forest and sequence-based models as a way to detect zero-day vulnerabilities in software systems, but previous results show that these techniques do not work as well on the high-dimensional and multi-source nature of telemetry data used with Kubernetes. The algorithm developed by Ahmad et al. used an isolation forest for anomaly detection and had less success detecting anomalies than with other algorithms because of these conditions. Additionally, deepLog models the sequence of log events using a recurrent neural network, which made it suitable for use with monolithic applications, but it is not capable of capturing the interdependencies between pods or changes in microservices rapidly. There are new multi-signal observability approaches that combine logs, metrics, and traces into a unified observation method, but these approaches also cannot comprehensively address the challenge of zero-day resilience without using temporal and relational modeling capabilities.

F. Research Gaps and Positioning

As observed in the literature analysis, many research gaps exist in Kubernetes security, such as insufficient real-time Explainability, poor Multi-source Fusion, unconsidered Temporal-Relational Modeling, insufficient Human-in-the-Loop Controls, and poor Generalization in Zero-day attacks. The above research gaps are summarized in Table 1. To address these research gaps, this research proposes a Temporal Graph Framework with a focus on Kubernetes, Multi-source Telemetry, TGNN Detection, Explainable AI, and Safety-Centric Containment.

Area	Gap	Addressed By
XAI for Security	Limited real-time explainability	SHAP + GNNExplainer integration
Zero-Day Detection	Poor multi-source	performance TGNN-eBPF ensemble
Operational Safety	No HITL controls	HITL approval + containment
Privacy & Cloud Security	No Kubernetes-ready privacy mechanisms	Pod-aware security graphs
AI for Kubernetes	Scalability untested	Temporal, pod-aware graphs
eBPF Observability	Not linked to ML/GNN systems	eBPF-TGNN data fusion
Log Analytics	Weak multi-source fusion	Unified logs + flows + syscalls

TGNN Methods	Not Kubernetes-specific; high overhead	Optimized TGNN with streaming inference
Adversarial Robustness	Performance degradation under topology changes	Temporal windowing and graph evolution modeling
Real-World Deployment	Limited validation on production clusters	Deployment-ready design with planned real-world dataset integration

Table 1. An overview of research gaps and our methodology

III. SYSTEM ARCHITECTURE

A. High-Level Design

Telemetry → Ingestion → Temporal Graph Builder → Graph Storage → ML Inference → XAI Layer → Alerts → Containment Operator

Fig. 1. High-Level Pipeline

B. Component Details

Telemetry Integration: eBPF-based monitoring with Cilium [3][4] is combined with audit logs and Istio traces.

C. Temporal Graph Construction Algorithm:

Algorithm 1: Temporal Graph Construction

Input: Telemetry events E , time window Δt

Output: Temporal graph sequence G_t

While new events are coming in:

- Read Event: Extract the time, source (Pod), and destination.
- Check Time: Determine the time "bucket" (Window) that this event is related to.
- Update Graph:
 - Add Nodes (the Pods/Processes).
 - Add Edges (the connection/flow between them).

When the Time Window is over (e.g., every 5 mins):

- Summarize: Gather the counts (traffic volume, execution rates).
- Save: Write out the graph and stats to a file (Parquet).
- Normalize: Adjust the figures for uniform analysis.
- Reset: Free up memory and prepare a new graph for the next window.

D. Anomaly Detection Pipeline

Baseline detectors include an Autoencoder with an MSE threshold, an LSTM-AE that uses a sequence length of 50, and DeepLog with a window size of 10 [12]. The graph-based model uses a two-layer GATConv TGNN [6] to [9]. Final decisions are made through weighted ensemble voting.

E. Explainability

The incorporation of explainability functions takes place within the detection system. The SHAP algorithm [10] delivers feature-level attribution of baseline model predictions via SHAP values. The GNNExplainer [7] illuminates which specific targeted subgraphs or patterns

contributed most significantly to a detection in the TGNN. Anomalies detected by the pipeline are correlated to the MITRE ATT&CK tactic to enhance the ability of the analyst to comprehend and interpret the analyzed data (i.e. using T1059 for elevated execve activity).

A Kubernetes operator creates a XML Schema (CRD) for performing containment functions. All methods of mitigation must be reviewed through a human-in-the-loop model and therefore require human validation to mitigate the risk of operational failure and to ensure that inadvertent disruption of service occurs during automated response actions.

F. TGNN Objective Function

The TGNN is optimized using an objective function with an anomaly-centric perspective, which incorporates both structural and temporal aspects:

$$L = \lambda_1 \cdot L_{rec} + \lambda_2 \cdot L_{temp}$$

here L_{rec} is a representation of the loss of reconstruction on node and edge features, while L_{temp} penalizes deviations in a series of time windows.

IV. IMPLEMENTATION AND EVALUATION

A. Experimental Setup

The evaluation environment consisted of a single-node Minikube cluster with 4 GB of RAM and 2 vCPUs. For generating benign workloads and zero-day attack traces, simulators were used.

The detection pipeline aggregates eBPF syscall traces, Kubernetes audit records, and service mesh Telemetries, which can all be natively obtained in production Kubernetes environments without any need for architecture modifications for deployment. Baseline Detectors: Autoencoder based on error thresholding for reconstruction error, LSTM Autoencoder for temporal Anomaly Detection, and DeepLog based on syscall sequence analysis [12].

Although this present assessment centers on synthetic workloads, this framework will soon be capable of processing actual Kubernetes metrics and attack traces based on CVE vulnerabilities, all described under Section V.

B. Results




Fig. 2. Detection rates across 100 test windows.

Figure 2 illustrates the temporal stability of detection performance under alternating normal and attack conditions (50 normal, 50 attack).

Model	Accuracy	Precision	Recall	F1	Latency
Autoencoder	94%	0.91	0.96	0.93	187ms
LSTM-AE	93%	0.88	0.95	0.91	212ms
DeepLog	91%	0.85	0.93	0.89	156ms
Ensemble	95%	0.92	0.96	0.94	198ms

Table 2. Detection performance across normal and attack workloads.

The ensemble model achieves the highest F1 score while maintaining low inference latency. AUC-PR will be reported in extended evaluation.




Fig. 3. Comparative performance of baseline and ensemble models.

Figure 3 plots the classification results of the individual models and the ensemble method. It clearly shows that the ensemble method achieves better accuracy (95%) and F1-score (0.94) compared to the individual models. This suggests that the temporal, sequential, and graph features enhance the robustness of the models.




Fig. 4: Inference Latency Over Test Windows

As shown in Fig. 4, inference latency remains consistently below operational thresholds.

Figure 3 summarizes comparative detection performance across models, while Figure 4 presents inference latency

behavior across test windows. The ensemble model has been seen to have the best trade-offs in terms of accuracy (95%), precision (92%), and recall of 96% while having inference latency of less than 500 ms in all the windows.

Cluster Size	Avg. Latency	Memory Usage	TGNN Load
Minikube	198 ms	Low	Low
5 Nodes	~310 ms	Medium	Medium
10 Nodes	~470 ms	High	High

Table 3. Projected Latency Trends Across Cluster Sizes

Table 3 shows projected scalability trends deduced from expected telemetry growth and TGNN message passing complexity, and is not measured benchmarks. Because multi-node clusters were not physically deployed during this phase, CPU utilization, memory growth, and inference latency are estimated on projected telemetry expansion, not from runtime profiled. Memory usage grows with cluster size due to the number of graph nodes and edges, and due to increased TGNN message-passing in higher inference time CPU utilization.

C. MITRE ATT&CK Mapping

Detected anomalies mapped to tactics:

- T1059: Command execution, execve bursts
- T1041: Exfiltration over C2, unusual outbound connections
- T1548: Privilege escalation, capability transitions
- T1570: Lateral movement, inter-pod reconnaissance

The interpretability of the anomalies was also investigated by relating the model outputs to the tactics of the MITRE ATT&CK framework via SHAP and GNNExplainer techniques. The following analysis is performed on the same set of detected anomalies produced during the evaluation phase described in Section IV-B.




Fig. 5: MITRE ATT&CK Tactic Distribution

Figure 5 represents how MITRE ATT&CK tactics are distributed based on detected anomalies. The highest number of detected behaviors accounts for command execution (T1059), followed by data exfiltration (T1041) and privilege escalation (T1548) tactics. What the distribution reveals is the presence of anomalies across various stages of the attacks, such as execution, privileged access, lateral movement, and exfiltration instead of single behaviors.




Fig. 6: Shap Feature Importance

For better explainability by analysts, SHAP and GNNExplainer were run on chosen anomaly data points. As shown in Figure 6, high calls for "execve", unusual output traffic, and high inter-pod communication were prominent feature contributors, which corresponded with the MITRE ATT&CK tactic "T1059: Command Execution."

D. Analyst Usability and Alert Interpretation

In a real operations setting, quality of detection by itself is not sufficient. Security analysts must quickly comprehend and react to alerts. The proposed framework displays each alert with its anomaly score, the corresponding MITRE ATT&CK tactic, and a short explanation derived from SHAP features and the graph context, including pod, process, and communication links. This is supposed to lessen the investigation time and also facilitate informed containment decisions.

A small pilot evaluation is set up with a group of security students or engineers. They will review sample alerts and rate them using a Likert scale based on clarity, actionability, usefulness of explanations, and their confidence in dealing with potential false positives. The planned pilot study is designed to provide information that will guide the future refinement of changes in alert formats and explanation details before their wide deployment in operational environments.

E. Comparison with Existing Security Tools and Models

Table 4 describes a qualitative comparative analysis of our framework with existing runtime security solutions and research models. Current runtime security solutions such as Falco and Sysdig mainly make use of rules or signature matching. They can detect known attacks but not zero-day attacks. Research models such as strGNN have capabilities in modeling over time and structure but aren't Kubernetes-native or explainable.

On the other hand, a proposed framework incorporates temporal graph modeling with awareness of Kubernetes, zero-day attack-capable anomaly analysis, explicable AI with MITRE ATT&CK mappings, and human-in-the-loop containment.

Tool/Model	Detection Type	Zero-Day Support	Explainability (XAI)	Kubernetes-Aware	Reported F1
Falco	Rule-Based	No	No	Yes	Not Reported

Sysdig	Signature-Based	No	No	Yes	Not Reported
strGNN	Research TGNN	Partial	No	No	≈ 0.88
Proposed	TGNN + XAI	Yes	Yes	Yes	0.94

Table 4. Qualitative Comparison to existing tools and models.

Reported F1 values are taken from the respective publications when available; several production tools do not report F1 metrics.

F. Risk Assessment and Human-in-the-Loop (HITL) Safety Analysis

Kernel-focused telemetry data acquisition using eBPF can be a problem when it sorts with personal information relating to process metadata. To address this problem, the solution does not inspect payload information but focuses on anonymous aggregated statistics.

Inference is conducted on graph summaries over a window rather than event data, thereby maintaining projected computational complexity at an acceptable level for real-time Kubernetes system analysis. While automation in both detection and containment remains very helpful in a quick response implementation, it can pose operational risks if not performed correctly. To address this challenge, a human-in-the-loop approach is used in this framework.

Table 5 below highlights the main operational risks considered and methods to mitigate them. A safety-first approach strikes a balance between automation and human intervention to ensure less possibility of service interruption but with a quick response time.

Risk	Cause	Impact	Mitigation
False positives	Traffic spikes or workload bursts	Unnecessary pod isolation	Analyst verification via HITL
Wrong containment	Sensor noise or feature drift	Service disruption	Manual override and rollback
Alert overload	Repeated anomaly triggers	Analyst fatigue	Alert deduplication and throttling

Table 5. Operational Risks and Human-in-the-Loop Safeguards

V. LIMITATIONS

- **Synthetic Evaluation:** Experimental evaluations were based on synthetic benign workloads and simulated zero-day attack scenarios, thus preventing any allowance of direct generalization to Kubernetes clusters in real-world applications.
- **Single-Node Testbed:** All experiments utilized a single-node Minikube deployment, thus providing no direct measurements of system behavior on production-type, larger scale clusters.
- **Projected Scalability:** Multi-node scalability results were based on expected trends that were not

empirically measured on distributed Kubernetes clusters.

- **Planned Usability Study:** The usability analyses to date are preliminary; a pilot human study will be conducted, but quantitative usability scores and response-time measurements will be provided in future publication(s).
- **Production Validation:** Validation of results via comparative analysis of results from several live production clusters and publicly available attack datasets are currently broken down into further tasks that will be addressed as future work.

Future work includes federated learning to enable privacy-preserving TGNN training across distributed Kubernetes clusters.

VI. CONCLUSION

A proof-of-concept solution for Zero-Day attacks in microservices with Kubernetes using TGNNs and Explainable AI is presented in this research work. The model represents Kubernetes workloads in a form of temporal graphs constructed from a wide variety of sources. Therefore, temporal inter-pod and intra-process relationships, which might not be easily detected using traditional model solutions, are incorporated in this model. The combination of TGNN with a baseline detector enhances both accuracy and latency in this model.

Explainability is incorporated in this work using SHAP and GNNExplainer to facilitate transparent analysis of alerts and their alignment with MITRE ATT&CK attack tactics. To promote operational safety, a human-in-the-loop containment strategy is used, which supports automatic responses without undermining service continuity. Experimental evaluation conducted in a Minikube environment with synthetic zero-day attacks shows excellent performance with a recall rate of 93–96% and an F1-score of 0.94.

The scope of this assessment will be confined to synthetic workloads and single node environments. Future research will include validating this framework with real-world Kubernetes Telemetries, attack datasets, and multi-node environments. In Summary, this approach emphasizes how Temporal Relational Modeling and Explainability, in combination with safety for proactive zero-day defense in cloud native systems, can prove beneficial.

REFERENCES

- [1] T. Arif, B. Jo and J. H. Park, "A comprehensive survey of privacy-enhancing and trust-centric cloud-native security techniques against cyber threats," *Sensors*, vol. 25, no. 8, p. 2350, 2025.
- [2] H. Pitkar, Enhancing Kubernetes security with AI: Anomaly detection for cloud-based workloads, 2025.
- [3] Hussain, A. Aziz, H. J. Syed and S. Raza, "Preventing IP spoofing in Kubernetes using eBPF," *Computers, Materials & Continua*, vol. 84, no. 2, p. 3105–3124, 2025.
- [4] P. Murugan, "Enhancing container security using eBPF-based detection mechanisms for real-time threat monitoring and system-level visibility," Zenodo, 2022.
- [5] N. A. Mohammed, "Cloud-native technologies AI-enhanced observability: Machine learning pipeline for Kubernetes log analytics in EKS environments," *World Journal of Advanced Research and Reviews*, vol. 26, no. 1, p. 888–896, 2025.
- [6] Ben-Yair, P. Rogovoy and N. Zaidenberg, "AI & eBPF-based performance anomaly detection system," in *Proc. 12th ACM International Conference on Systems and Storage (SYSTOR '19)*, Haifa, 2019.
- [7] L. Cai, Z. Chen, C. Luo, J. Gui, J. Ni, D. Li and H. Chen, "Structural temporal graph neural networks for anomaly detection in dynamic graphs," in *Proc. 30th ACM International Conference on Information and Knowledge Management (CIKM '21)*, 2021.
- [8] Dutt, D. Jang, J. Nadkarni, K. Su and A. Gandhi, GUIDE-GNN based unified incident detection for microservices application deployments
- [9] Y. Fu, C. Zhou, J. Li and L. Chen, "Long-term evolutionary patterns matter: Self-supervised anomaly detection on dynamic graphs," *Knowledge-Based Systems*, vol. 311, p. 113049, 2025.
- [10] AIIN CYBERSECURITY, A comprehensive survey on explainable AI in cybersecurity domain, Society for Electronic Transactions and Security (SETS).
- [11] R. Alenezi and S. A. Ludwig, "Explainability of cybersecurity threats data using SHAP," in *Proc. 2021 IEEE Symposium Series on Computational Intelligence (SSCI)*, 2021.
- [12] P. T. Thanusma, A. Srinivas, M. Jayasurya, S. Ashwini, S. M. Vali, T. A. Sri and D. Santhadevi, "Anomaly detection for zero-day exploit identification in network traffic using isolation forest," *International Journal of Advances in Computer Science and Cloud Computing (IJACSCC)*, vol. 12, no. 2, p. 27–34, 2024.
- [13] M. Ahmad, K. Imran, H. Attar and A. Solyman, "Isolation forest-based anomaly detection through generation offers and optimal dispatch data," in *Proc. 3rd International Conference on Emerging Power Technologies*, 2025.
- [14] R. Guntupalli, "AI-driven anomaly detection and root cause analysis: Using machine learning on logs, metrics, and traces to detect subtle performance anomalies, security threats, or failures in complex cloud environments," *World Journal of Advanced Research and Reviews*, vol. 26, no. 2, p. 874–879, 2025
- [15] M. Du, F. Li, G. Zheng, and V. Srikumar, "DeepLog: Anomaly detection and diagnosis from system logs through deep learning," in *Proc. ACM SIGSAC Conf. on Computer and Communications Security (CCS)*, Dallas, TX, USA, Oct. 2017.
- [16] V. Govindarajan and J. H. Muzamal, "A transformer-based framework for network intrusion detection using graph learning and contrastive optimization," *Scientific Reports*, vol. 15, p. 20511, 2025.
- [17] A. Jakkaraju, "Graph Neural Networks for Anomaly Detection in Cloud Infrastructure," *Journal of Computer and Communications*, vol. 13, no. 10, pp. 102–116, Oct. 2025.
- [18] Z. Zhang, Z. Zhu, C. Xu, J. Zhang, and S. Xu, "Towards accurate anomaly detection for cloud system via graph-enhanced contrastive learning," *Complex & Intelligent Systems*, vol. 11, no. 23, 2025.